

After working the code we can start to see the shape that this data take, specifically who is in the data and who is missing. Since Grad Café submissions is a voluntary selection rather than sampled, invited, or followed up; the database is a self-selected record of postings rather than a representative sample of graduate applicants. The data show that it overrepresents research doctorates and prominent institutions: PhD results make up 70.85% of the 30,503 entries, while just 50 of the 1,240 universities account for 60.54% of all entries, even though master's applicants nationally far outnumber PhD applicants. Posting are also concentrated around decisions, such as 77.41% of entries were added in February and March 2026 and overwhelmingly in Fall 2026. These three months represent 98.57% of the data. Consequently, Question 5's Fall 2025 acceptance rate of 48.22% is based on only 197 late-posted entries from an older cycle, and is not a random cross-section of Fall 2025 applicants. Lastly, the site provides no applicant identifiers, so we have no way to telling it this is a one-to-one person to entry. Meaning repeats and duplicate submissions cannot be reliably detected, and "entries" must not be interpreted as "applicants."

The reported value also cannot be read at face value, because disclosure is sparse, uneven, and unverified. GRE Quantitative appears in only 7.95% of entries; among scores on the valid scale, the mean is 165.82, 72.08% are at least 165, and 31.95% are reported as a perfect 170. Compared to the assignment's roughly 157 population benchmark, this is strong evidence that high scorers disproportionately choose to report. Missingness also varies with the outcomes: accepted applicants report a GPA 65.80% of the time, compared with the 50.90% for waitlisted applicants. Entries omitting nationality have a 55.08% acceptance rate versus 38.49% for American and 34.55% for International entries. Data-quality problems reinforce the concern: 1,427 of 2,426 GRE entries are actually 260–340 composite totals entered in a section-score field, 135 reported GPAs exceed 4.0, and some submissions are jokes/trolling or fabricated universities. The valid-scale rule is therefore consequential rather than neutral. The stored GRE Quantitative mean is 259.96, but the filtered mean is 165.82. The range restriction further obscures relationships: accepted applicants average a 3.76 GPA and rejected applicants 3.78, while 88.71% of reported GPAs are at least 3.5 and 14.57% are exactly 4.0. This ceiling effect in a self-selected pool does not show that GPA is irrelevant to admission. In addition, university and program names were standardized with a small local language model and fuzzy matching, so Question 11's ranking depends on that mapping and still separates some variants, such as "Yale" and "Yale University." Taken together, the clearest cautions are the elevated GRE mean relative to the roughly 157 benchmark, outcome-dependent GPA reporting, and the very small Fall 2025 sample: each demonstrates the gap between what the database

contains and what it can support as a general claim.

Figure	Value
Entries / distinct universities	30,503 / 1,240
GRE Quantitative reported	7.95%
Mean valid GRE Q / mean as stored	165.82 / 259.96
GPA reported: accepted / waitlisted	65.80% / 50.90%
Mean GPA: accepted / rejected	3.76 / 3.78
Fall 2025 entries behind Q5's 48.22%	197